

```
In [1]: import tensorflow as tf
import os
import pandas as pd
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, TensorBoard
from tensorflow.keras.losses import MeanSquaredError
from tensorflow.keras.metrics import RootMeanSquaredError
from tensorflow.keras.optimizers import Adam
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
In [9]: import datetime

# Define the path to store the logs
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True)
```

```
In [3]: zip_path = tf.keras.utils.get_file(
    origin='https://storage.googleapis.com/tensorflow/tf-keras-datasets/jena_climate_2009_2016.csv.zip',
    fname='jena_climate_2009_2016.csv.zip',
    extract=True)
csv_path, _ = os.path.splitext(zip_path)
```

```
In [10]: df = pd.read_csv(csv_path)
```

```
In [16]: df = df[5::6]
```

```
In [17]: df.index = pd.to_datetime(df['Date Time'], format='%d.%m.%Y %H:%M:%S')
```

```
In [18]: temp = df['T (degC)']
```

```
In [19]: # [[[1], [2], [3], [4], [5]]] [6]
# [[[2], [3], [4], [5], [6]]] [7]
# [[[3], [4], [5], [6], [7]]] [8]

def df_to_X_y(df, window_size=5):
    df_as_np = df.to_numpy()
    X = []
    y = []
    for i in range(len(df_as_np)-window_size):
        row = [[a] for a in df_as_np[i:i+window_size]]
        X.append(row)
        target = df_as_np[i+window_size]
        y.append(target)
    return np.array(X), np.array(y)
```

```
In [20]: WINDOW_SIZE = 5
X1, y1 = df_to_X_y(temp, WINDOW_SIZE)
X1.shape, y1.shape
```

```
Out[20]: ((70086, 5, 1), (70086,))
```

```
In [21]: def split_data(X, y, train_percent, val_percent_of_train):
    total_samples = len(X)
    train_size = int(total_samples * train_percent)
    val_size = int(train_size * val_percent_of_train)

    # Recompute train_size to exclude the validation set from the original train set
    train_size -= val_size

    X_train, y_train = X[:train_size], y[:train_size]
    X_val, y_val = X[train_size:train_size + val_size], y[train_size:train_size + val_size]
    X_test, y_test = X[train_size + val_size:], y[train_size + val_size:]

    return (X_train, y_train), (X_val, y_val), (X_test, y_test)

# Example usage:
train_percent = 0.7 # 70% of data is initially considered for training
val_percent_of_train = 0.15 # 15% of the initial training set is for validation

(X_train, y_train), (X_val, y_val), (X_test, y_test) = split_data(X1, y1, train_percent, val_percent_of_train)
```

```
In [25]: model3 = Sequential()
model3.add(InputLayer((5, 1)))
model3.add(GRU(64))
model3.add(Dense(8, 'relu')) #Intermediate Dense Layer-The first Dense layer reduces the dimension from the LSTM's out
model3.add(Dense(1, 'linear')) #Output Dense Layer This layer is crucial as it maps the processed features from the pr

model3.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| gru_1 (GRU) | (None, 64) | 12,864 |
| dense_2 (Dense) | (None, 8) | 520 |
| dense_3 (Dense) | (None, 1) | 9 |

Total params: 13,393 (52.32 KB)

Trainable params: 13,393 (52.32 KB)


Non-trainable params: 0 (0.00 B)

```
In [26]: cp1 = ModelCheckpoint('model3/model_checkpoint.keras', save_best_only=True)
es1 = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```


```
In [27]: model3.compile(loss=MeanSquaredError(), optimizer=Adam(learning_rate=0.0001), metrics=[RootMeanSquaredError()])
```

```
In [28]: model3.fit(
          x=X_train,
          y=y_train,
          validation_data=(X_val, y_val),
          epochs=10,
          callbacks=[cp1, es1, tensorboard_callback]
        )
```


Epoch 1/10

1304/1304  5s 3ms/step - loss: 90.2006 - root_mean_squared_error: 9.3374 - val_loss: 8.1104 - val_root_mean_squared_error: 2.8479


Epoch 2/10

1304/1304  4s 3ms/step - loss: 6.3038 - root_mean_squared_error: 2.4940 - val_loss: 1.9036 - val_root_mean_squared_error: 1.3797


Epoch 3/10

1304/1304  3s 3ms/step - loss: 1.7224 - root_mean_squared_error: 1.3102 - val_loss: 1.0137 - val_root_mean_squared_error: 1.0068


Epoch 4/10

1304/1304  3s 3ms/step - loss: 0.9584 - root_mean_squared_error: 0.9785 - val_loss: 0.7813 - val_root_mean_squared_error: 0.8839


Epoch 5/10

1304/1304  3s 3ms/step - loss: 0.7535 - root_mean_squared_error: 0.8678 - val_loss: 0.7014 - val_root_mean_squared_error: 0.8375


Epoch 6/10

1304/1304  3s 3ms/step - loss: 0.6878 - root_mean_squared_error: 0.8293 - val_loss: 0.6757 - val_root_mean_squared_error: 0.8220


Epoch 7/10

1304/1304  3s 3ms/step - loss: 0.6965 - root_mean_squared_error: 0.8338 - val_loss: 0.6751 - val_root_mean_squared_error: 0.8216


Epoch 8/10

1304/1304  3s 3ms/step - loss: 0.6294 - root_mean_squared_error: 0.7931 - val_loss: 0.6561 - val_root_mean_squared_error: 0.8100

Epoch 9/10

1304/1304  3s 3ms/step - loss: 0.6298 - root_mean_squared_error: 0.7935 - val_loss: 0.6616 - val_root_mean_squared_error: 0.8134

Epoch 10/10

1304/1304  3s 3ms/step - loss: 0.6494 - root_mean_squared_error: 0.8058 - val_loss: 0.6622 - val_root_mean_squared_error: 0.8137

Out[28]: <keras.src.callbacks.history.History at 0x2dc81983150>

```
In [29]: from tensorflow.keras.models import load_model  
model3 = load_model('model3/model_checkpoint.keras')
```

```
In [30]: train_predictions = model3.predict(X_train).flatten()
```

1304/1304 ————— 2s 1ms/step

```
In [31]: train_results = pd.DataFrame(data={'Train Predictions':train_predictions, 'Actuals':y_train})
```



```
In [32]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Calculate residuals
train_results['Residuals'] = train_results['Actuals'] - train_results['Train Predictions']

# Calculate error metrics
mae = mean_absolute_error(train_results['Actuals'], train_results['Train Predictions'])
mse = mean_squared_error(train_results['Actuals'], train_results['Train Predictions'])
rmse = np.sqrt(mse)

# Metrics for display
metrics_text = f"MAE: {mae:.2f}\nMSE: {mse:.2f}\nRMSE: {rmse:.2f}"


# Plot Predictions vs. Actuals
plt.figure(figsize=(12, 5))
plt.plot(train_results['Train Predictions'], label='Train Predictions')
plt.plot(train_results['Actuals'], label='Actuals')
plt.title('Predicted vs. Actual Values')
plt.ylabel('Values')
plt.legend()
plt.text(0.01, 0.99, metrics_text, verticalalignment='top', horizontalalignment='left', transform=plt.gca().transAxes,
plt.show())


# Residual Plot for Magnitude-Residual Relationship
# Calculate squared residuals
train_results['Squared Residuals'] = (train_results['Actuals'] - train_results['Train Predictions'])**2
plt.figure(figsize=(10, 6))
plt.scatter(train_results['Actuals'], train_results['Squared Residuals'], alpha=0.5)
plt.title('Magnitude-Residual Relationship')
plt.xlabel('Actual Values')
plt.ylabel('Squared Residuals')
plt.grid(True)
plt.text(0.01, 0.99, metrics_text, verticalalignment='top', horizontalalignment='left', transform=plt.gca().transAxes,
plt.show())


# Histogram of Residuals
plt.figure(figsize=(8, 6))
plt.hist(train_results['Residuals'], bins=20, edgecolor='black', alpha=0.7)
```



```
plt.title('Histogram of Residuals')  
plt.xlabel('Residuals')  
plt.ylabel('Frequency')  
plt.text(0.99, 0.99, metrics_text, verticalalignment='top', horizontalalignment='right', transform=plt.gca().transAxes)  
plt.show()
```







```
In [33]: import matplotlib.pyplot as plt

# Set the figure size
plt.figure(figsize=(12, 5))


# Plot training predictions
plt.plot(train_results['Train Predictions'][:200], label='Train Predictions')

# Plot actual values
plt.plot(train_results['Actuals'][:200], label='Actuals')

# Adding legend to the plot
plt.legend()

# Adding titles and labels
plt.title('Comparison of Training Predictions and Actuals')
plt.xlabel('Sample Index')
plt.ylabel('Values')

# Display the plot
plt.show()
```



```
In [34]: test_predictions = model3.predict(X_test).flatten()
```

658/658 ————— 1s 997us/step


```
In [35]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Assuming you have already loaded the model and have test_predictions and y_test available
# Create a DataFrame to hold test predictions and actual values
test_results = pd.DataFrame({
    'Test Predictions': test_predictions,
    'Actuals': y_test
})

# Calculate residuals
test_results['Residuals'] = test_results['Actuals'] - test_results['Test Predictions']

# Calculate error metrics
test_mae = mean_absolute_error(test_results['Actuals'], test_results['Test Predictions'])
test_mse = mean_squared_error(test_results['Actuals'], test_results['Test Predictions'])
test_rmse = np.sqrt(test_mse)


# Metrics for display
test_metrics_text = f"MAE: {test_mae:.2f}\nMSE: {test_mse:.2f}\nRMSE: {test_rmse:.2f}"


# Plot Predictions vs. Actuals for Test Data
plt.figure(figsize=(12, 5))
plt.plot(test_results['Test Predictions'], label='Test Predictions')
plt.plot(test_results['Actuals'], label='Actuals')
plt.title('Test Predicted vs. Actual Values')
plt.ylabel('Values')
plt.legend()
plt.text(0.01, 0.99, test_metrics_text, verticalalignment='top', horizontalalignment='left', transform=plt.gca().trans)
plt.show()

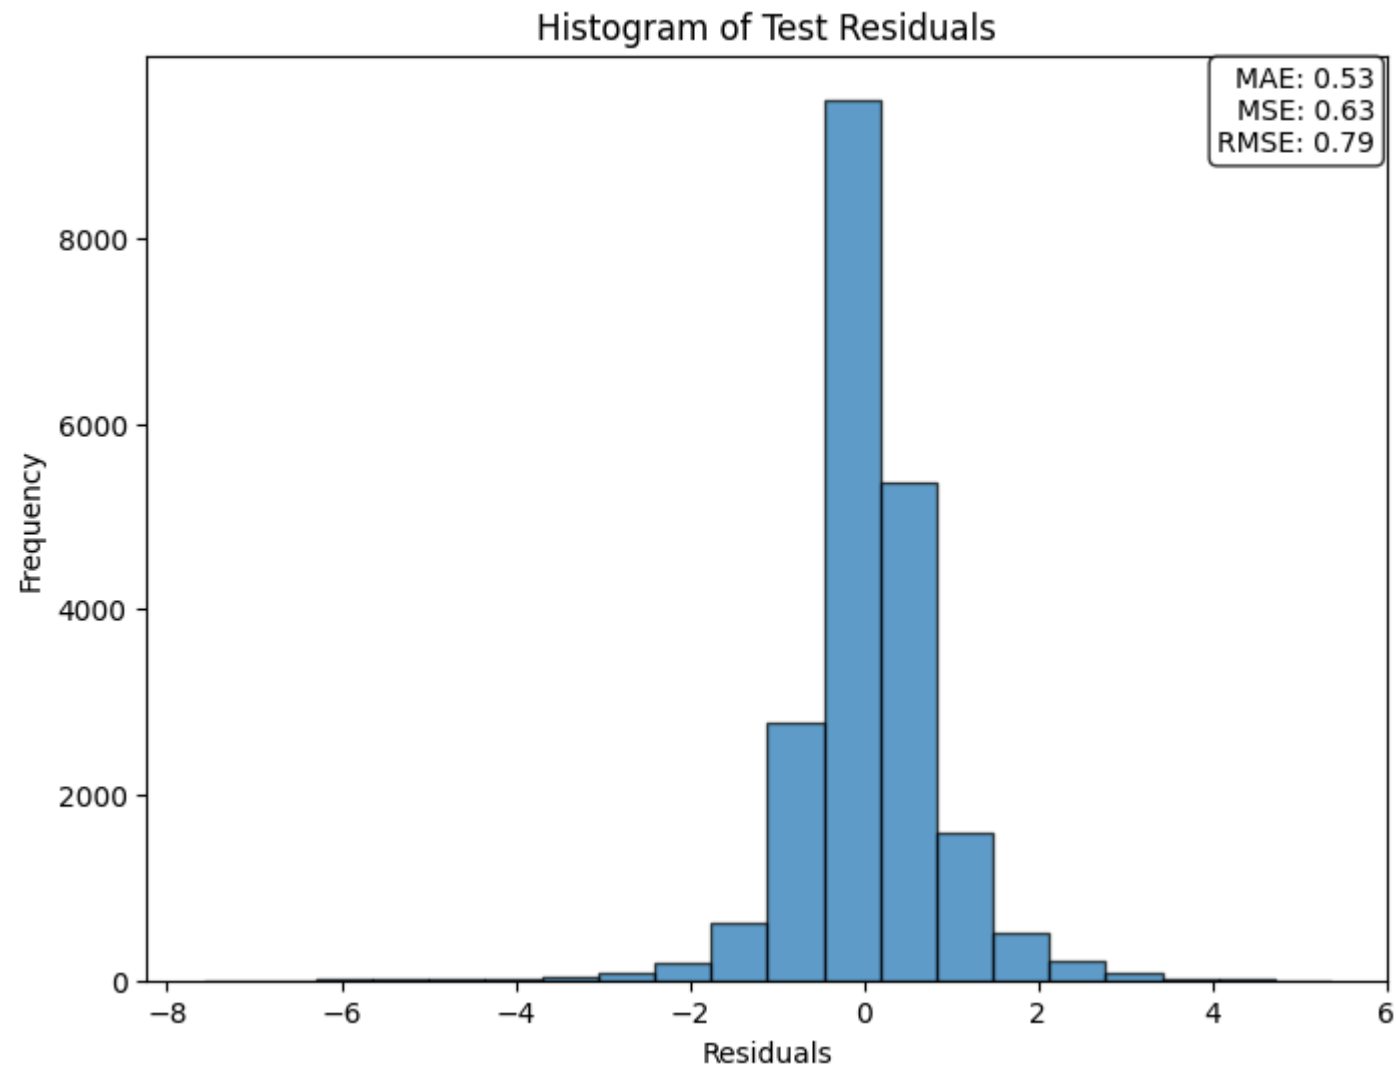
# Residual Plot for Magnitude-Residual Relationship for Test Data
# Calculate squared residuals
test_results['Squared Residuals'] = (test_results['Actuals'] - test_results['Test Predictions'])**2
plt.figure(figsize=(10, 6))
plt.scatter(test_results['Actuals'], test_results['Squared Residuals'], alpha=0.5)
plt.title('Test Magnitude-Residual Relationship')
plt.xlabel('Actual Values')
plt.ylabel('Squared Residuals')
```

```
plt.grid(True)
plt.text(0.01, 0.99, test_metrics_text, verticalalignment='top', horizontalalignment='left', transform=plt.gca().trans
plt.show()

# Histogram of Residuals for Test Data
plt.figure(figsize=(8, 6))
plt.hist(test_results['Residuals'], bins=20, edgecolor='black', alpha=0.7)
plt.title('Histogram of Test Residuals')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.text(0.99, 0.99, test_metrics_text, verticalalignment='top', horizontalalignment='right', transform=plt.gca().tran
plt.show()
```







```
In [36]: import matplotlib.pyplot as plt

# Set the figure size
plt.figure(figsize=(12, 5))

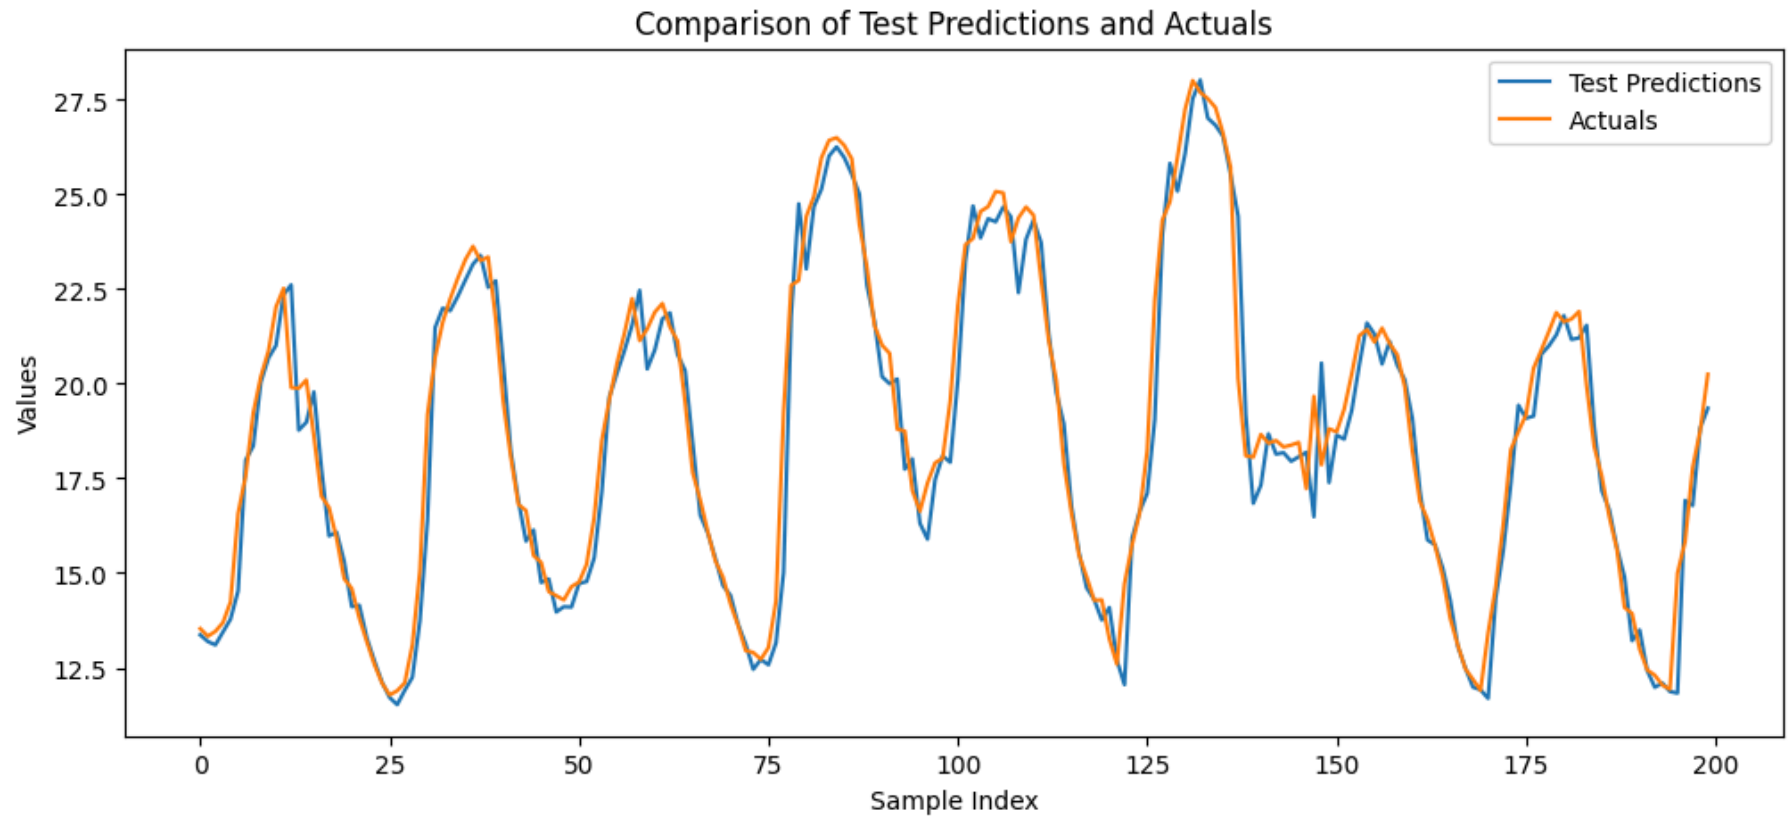
# Plot test predictions for the first 200 entries
plt.plot(test_results['Test Predictions'][:200], label='Test Predictions')

# Plot actual values for the first 200 entries
plt.plot(test_results['Actuals'][:200], label='Actuals')

# Adding legend to the plot
plt.legend()

# Adding titles and labels
plt.title('Comparison of Test Predictions and Actuals')
plt.xlabel('Sample Index')
plt.ylabel('Values')

# Display the plot
plt.show()
```



In []: